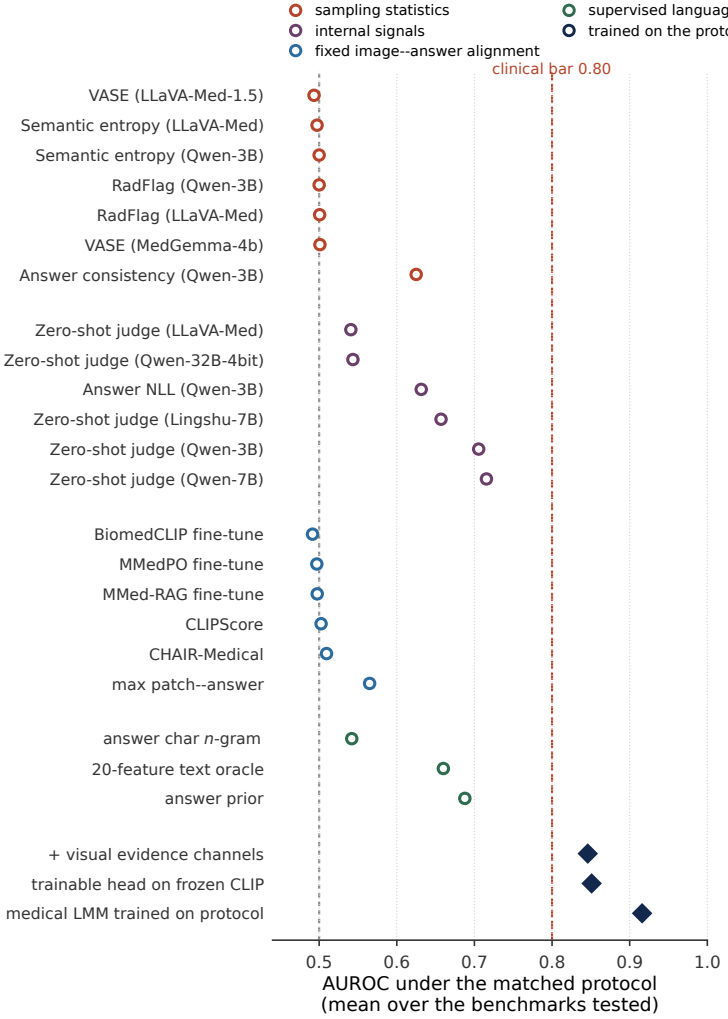
sampling statistics
internal signals
fixed image--answer alignment
supervised language statistics
trained on the protocol
clinical bar 0.80
VASE (LLaVA-Med-1.5)
Semantic entropy (LLaVA-Med)
Semantic entropy (Qwen-3B)
RadFlag (Qwen-3B)
RadFlag (LLaVA-Med)
VASE (MedGemma-4b)
Answer consistency (Qwen-3B)
Zero-shot judge (LLaVA-Med)
Zero-shot judge (Qwen-32B-4bit)
Answer NLL (Qwen-3B)
Zero-shot judge (Lingshu-7B)
Zero-shot judge (Qwen-3B)
Zero-shot judge (Qwen-7B)
BiomedCLIP fine-tune
MMedPO fine-tune
MMed-RAG fine-tune
CLIPScore
CHAIR-Medical
max patch--answer
answer char n-gram
20-feature text oracle
answer prior
+ visual evidence channels
trainable head on frozen CLIP
medical LMM trained on protocol
0.5
0.6
0.7
0.8
0.9
1.0
AUROC under the matched protocol
(mean over the benchmarks tested)